

CSI 5155 Machine Learning

Assignment 1 – Report

Lecturer: Herna Viktor (hviktor@uottawa.ca)

Student: Kelvin Mock (kmock073@uOttawa.ca)

Student ID: 300453668

GitHub: <https://github.com/kmock930/Drug-Consumption-Machine-Learning-analysis.git>

Contents

Introduction	3
Models Evaluation	3
Classification without sampling	3
Classification with the Under-sampling method	5
Classification with the Over-sampling method	7
Classification with a combination of sampling methods	9
Further Analysis	11
Lesson Learnt.....	12
Processes in the Pipeline	12
Analysis from the Evaluation Metrics	12
Issue in the Datasets	12
Addressing the Issue	12
Conclusion	13

Introduction

In the given classification problem, I applied a machine learning pipeline with 6 different classifiers on 2 datasets respectively based on the following approaches: classifying with fine-tuned parameters (from the *Random Search* algorithm), classifying with under-sampled data, classifying with over-sampled data, and classifying with a combination of data from both sampling methods.

In this report, I am going to show their performances with *confusion matrices*, some insightful metrics (including the *recall* and the *precision*) of each classifier, as well as *Receiver-Operating Characteristic curves (ROCs)* which show the performances of each classifier in terms of an *Area Under Curve (AUC)*.

Models Evaluation

Classification without sampling

Chocolate Dataset

Decision Tree classifier

Predicted	user	non-user	All
Actual			
user	621.0	0.0	621.0
non-user	8.0	0.0	8.0
All	629.0	0.0	629.0

Precision: 0.9872813990461049

Recall: 1.0

Random Forest classifier

Predicted	user	non-user	All
Actual			
user	621.0	0.0	621.0
non-user	8.0	0.0	8.0
All	629.0	0.0	629.0

Precision: 0.9872813990461049

Recall: 1.0

SVM classifier

Predicted	user	non-user	All
Actual			
user	621.0	0.0	621.0
non-user	8.0	0.0	8.0
All	629.0	0.0	629.0

Precision: 0.9872813990461049

Recall: 1.0

Gradient Boosting classifier

Predicted	user	non-user	All
Actual			
user	618	3	621
non-user	8	0	8
All	626	3	629

Precision: 0.987220447284345

Recall: 0.9951690821256038

MLP classifier

Predicted	user	non-user	All
Actual			
user	617	4	621
non-user	8	0	8
All	625	4	629

Precision: 0.9872

Recall: 0.9935587761674718

KNN classifier

Predicted	user	non-user	All
Actual			
user	621.0	0.0	621.0
non-user	8.0	0.0	8.0
All	629.0	0.0	629.0

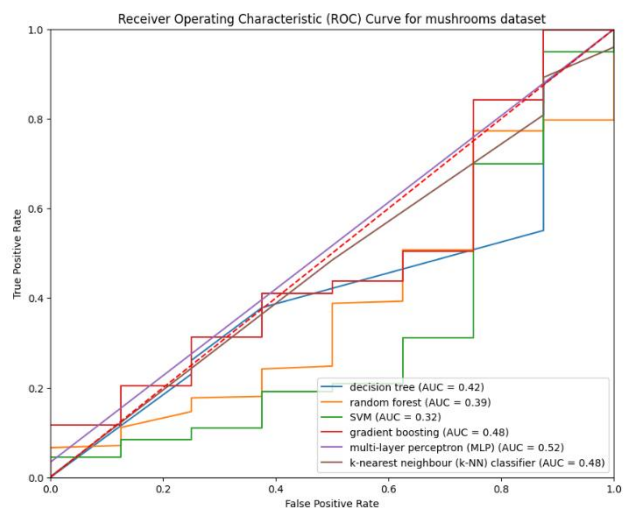

Precision: 0.9872813990461049

Recall: 1.0

Mushrooms Dataset

<div><div>Decision Tree classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>170</td><td>451</td><td>621</td></tr><tr><td>non-user</td><td>2</td><td>6</td><td>8</td></tr><tr><td>All</td><td>172</td><td>457</td><td>629</td></tr></table><div>Precision: 0.9883720930232558 Recall: 0.9883720930232558</div></div>	Predicted	user	non-user	All	Actual				user	170	451	621	non-user	2	6	8	All	172	457	629	<div><div>Random Forest classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>221</td><td>400</td><td>621</td></tr><tr><td>non-user</td><td>4</td><td>4</td><td>8</td></tr><tr><td>All</td><td>225</td><td>404</td><td>629</td></tr></table><div>Precision: 0.9822222222222222 Recall: 0.355877616747182</div></div>	Predicted	user	non-user	All	Actual				user	221	400	621	non-user	4	4	8	All	225	404	629
Predicted	user	non-user	All																																						
Actual																																									
user	170	451	621																																						
non-user	2	6	8																																						
All	172	457	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	221	400	621																																						
non-user	4	4	8																																						
All	225	404	629																																						
<div><div>SVM classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>236</td><td>385</td><td>621</td></tr><tr><td>non-user</td><td>4</td><td>4</td><td>8</td></tr><tr><td>All</td><td>240</td><td>389</td><td>629</td></tr></table><div>Precision: 0.9833333333333333 Recall: 0.38003220611916266</div></div>	Predicted	user	non-user	All	Actual				user	236	385	621	non-user	4	4	8	All	240	389	629	<div><div>Gradient Boosting classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>211</td><td>410</td><td>621</td></tr><tr><td>non-user</td><td>5</td><td>3</td><td>8</td></tr><tr><td>All</td><td>216</td><td>413</td><td>629</td></tr></table><div>Precision: 0.9768518518518519 Recall: 0.3397745571658615</div></div>	Predicted	user	non-user	All	Actual				user	211	410	621	non-user	5	3	8	All	216	413	629
Predicted	user	non-user	All																																						
Actual																																									
user	236	385	621																																						
non-user	4	4	8																																						
All	240	389	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	211	410	621																																						
non-user	5	3	8																																						
All	216	413	629																																						
<div><div>MLP classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>222</td><td>399</td><td>621</td></tr><tr><td>non-user</td><td>4</td><td>4</td><td>8</td></tr><tr><td>All</td><td>226</td><td>403</td><td>629</td></tr></table><div>Precision: 0.9823008849557522 Recall: 0.357487922705314</div></div>	Predicted	user	non-user	All	Actual				user	222	399	621	non-user	4	4	8	All	226	403	629	<div><div>KNN classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>239</td><td>382</td><td>621</td></tr><tr><td>non-user</td><td>4</td><td>4</td><td>8</td></tr><tr><td>All</td><td>243</td><td>386</td><td>629</td></tr></table><div>Precision: 0.9835390946502057 Recall: 0.38486312399355876</div></div>	Predicted	user	non-user	All	Actual				user	239	382	621	non-user	4	4	8	All	243	386	629
Predicted	user	non-user	All																																						
Actual																																									
user	222	399	621																																						
non-user	4	4	8																																						
All	226	403	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	239	382	621																																						
non-user	4	4	8																																						
All	243	386	629																																						

ROC curve as a Summary



Classification with the Under-sampling method

Chocolate Dataset

<div><div>Decision Tree classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>589</td><td>32</td><td>621</td></tr><tr><td>non-user</td><td>8</td><td>0</td><td>8</td></tr><tr><td>All</td><td>597</td><td>32</td><td>629</td></tr></table><div>Precision: 0.9865996649916248 Recall: 0.9484702093397746</div></div>	Predicted	user	non-user	All	Actual				user	589	32	621	non-user	8	0	8	All	597	32	629	<div><div>Random Forest classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>481</td><td>140</td><td>621</td></tr><tr><td>non-user</td><td>4</td><td>4</td><td>8</td></tr><tr><td>All</td><td>485</td><td>144</td><td>629</td></tr></table><div>Precision: 0.9917525773195877 Recall: 0.7745571658615137</div></div>	Predicted	user	non-user	All	Actual				user	481	140	621	non-user	4	4	8	All	485	144	629
Predicted	user	non-user	All																																						
Actual																																									
user	589	32	621																																						
non-user	8	0	8																																						
All	597	32	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	481	140	621																																						
non-user	4	4	8																																						
All	485	144	629																																						
<div><div>SVM classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>553</td><td>68</td><td>621</td></tr><tr><td>non-user</td><td>8</td><td>0</td><td>8</td></tr><tr><td>All</td><td>561</td><td>68</td><td>629</td></tr></table><div>Precision: 0.9857397504456328 Recall: 0.8904991948470209</div></div>	Predicted	user	non-user	All	Actual				user	553	68	621	non-user	8	0	8	All	561	68	629	<div><div>Gradient Boosting classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>472</td><td>149</td><td>621</td></tr><tr><td>non-user</td><td>5</td><td>3</td><td>8</td></tr><tr><td>All</td><td>477</td><td>152</td><td>629</td></tr></table><div>Precision: 0.989517819706499 Recall: 0.7600644122383253</div></div>	Predicted	user	non-user	All	Actual				user	472	149	621	non-user	5	3	8	All	477	152	629
Predicted	user	non-user	All																																						
Actual																																									
user	553	68	621																																						
non-user	8	0	8																																						
All	561	68	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	472	149	621																																						
non-user	5	3	8																																						
All	477	152	629																																						
<div><div>MLP classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>493</td><td>128</td><td>621</td></tr><tr><td>non-user</td><td>7</td><td>1</td><td>8</td></tr><tr><td>All</td><td>500</td><td>129</td><td>629</td></tr></table><div>Precision: 0.986 Recall: 0.7938808373590982</div></div>	Predicted	user	non-user	All	Actual				user	493	128	621	non-user	7	1	8	All	500	129	629	<div><div>KNN classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>568</td><td>53</td><td>621</td></tr><tr><td>non-user</td><td>8</td><td>0</td><td>8</td></tr><tr><td>All</td><td>576</td><td>53</td><td>629</td></tr></table><div>Precision: 0.9861111111111112 Recall: 0.9146537842190016</div></div>	Predicted	user	non-user	All	Actual				user	568	53	621	non-user	8	0	8	All	576	53	629
Predicted	user	non-user	All																																						
Actual																																									
user	493	128	621																																						
non-user	7	1	8																																						
All	500	129	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	568	53	621																																						
non-user	8	0	8																																						
All	576	53	629																																						

Classification with the Over-sampling method

Chocolate Dataset

Decision Tree classifier

Predicted	user	non-user	All
Actual			
user	456	165	621
non-user	7	1	8
All	463	166	629

Precision: 0.9848812095032398
Recall: 0.7342995169082126

Random Forest classifier

Predicted	user	non-user	All
Actual			
user	427	194	621
non-user	6	2	8
All	433	196	629

Precision: 0.9861431870669746
Recall: 0.6876006441223832

SVM classifier

Predicted	user	non-user	All
Actual			
user	481	140	621
non-user	7	1	8
All	488	141	629

Precision: 0.985655737704918
Recall: 0.7745571658615137

Gradient Boosting

Predicted	user	non-user	All
Actual			
user	415	206	621
non-user	6	2	8
All	421	208	629

Precision: 0.9857482185273159
Recall: 0.6682769726247987

MLP classifier

Predicted	user	non-user	All
Actual			
user	466	155	621
non-user	5	3	8
All	471	158	629

Precision: 0.9893842887473461
Recall: 0.750402576489533

KNN classifier


Predicted	user	non-user	All
Actual			
user	421	200	621
non-user	4	4	8
All	425	204	629

Precision: 0.9905882352941177
Recall: 0.677938808373591

Mushrooms Dataset

<div><div>Decision Tree classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>431</td><td>190</td><td>621</td></tr><tr><td>non-user</td><td>4</td><td>4</td><td>8</td></tr><tr><td>All</td><td>435</td><td>194</td><td>629</td></tr></table><div>Precision: 0.9908045977011494 Recall: 0.6940418679549114</div></div>	Predicted	user	non-user	All	Actual				user	431	190	621	non-user	4	4	8	All	435	194	629	<div><div>Random Forest classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>428</td><td>193</td><td>621</td></tr><tr><td>non-user</td><td>5</td><td>3</td><td>8</td></tr><tr><td>All</td><td>433</td><td>196</td><td>629</td></tr></table><div>Precision: 0.9884526558891455 Recall: 0.6892109500805152</div></div>	Predicted	user	non-user	All	Actual				user	428	193	621	non-user	5	3	8	All	433	196	629
Predicted	user	non-user	All																																						
Actual																																									
user	431	190	621																																						
non-user	4	4	8																																						
All	435	194	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	428	193	621																																						
non-user	5	3	8																																						
All	433	196	629																																						
<div><div>SVM classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>481</td><td>140</td><td>621</td></tr><tr><td>non-user</td><td>7</td><td>1</td><td>8</td></tr><tr><td>All</td><td>488</td><td>141</td><td>629</td></tr></table><div>Precision: 0.985655737704918 Recall: 0.7745571658615137</div></div>	Predicted	user	non-user	All	Actual				user	481	140	621	non-user	7	1	8	All	488	141	629	<div><div>Gradient Boosting classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>411</td><td>210</td><td>621</td></tr><tr><td>non-user</td><td>6</td><td>2</td><td>8</td></tr><tr><td>All</td><td>417</td><td>212</td><td>629</td></tr></table><div>Precision: 0.9856115107913669 Recall: 0.6618357487922706</div></div>	Predicted	user	non-user	All	Actual				user	411	210	621	non-user	6	2	8	All	417	212	629
Predicted	user	non-user	All																																						
Actual																																									
user	481	140	621																																						
non-user	7	1	8																																						
All	488	141	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	411	210	621																																						
non-user	6	2	8																																						
All	417	212	629																																						
<div><div>MLP classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>467</td><td>154</td><td>621</td></tr><tr><td>non-user</td><td>5</td><td>3</td><td>8</td></tr><tr><td>All</td><td>472</td><td>157</td><td>629</td></tr></table><div>Precision: 0.989406779661017 Recall: 0.7520128824476651</div></div>	Predicted	user	non-user	All	Actual				user	467	154	621	non-user	5	3	8	All	472	157	629	<div><div>KNN classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>421</td><td>200</td><td>621</td></tr><tr><td>non-user</td><td>4</td><td>4</td><td>8</td></tr><tr><td>All</td><td>425</td><td>204</td><td>629</td></tr></table><div>Precision: 0.9905882352941177 Recall: 0.677938808373591</div></div>	Predicted	user	non-user	All	Actual				user	421	200	621	non-user	4	4	8	All	425	204	629
Predicted	user	non-user	All																																						
Actual																																									
user	467	154	621																																						
non-user	5	3	8																																						
All	472	157	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	421	200	621																																						
non-user	4	4	8																																						
All	425	204	629																																						

ROC curve as a Summary



Classification with a combination of sampling methods

Chocolate Dataset

<div><div>Decision Tree classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>319</td><td>302</td><td>621</td></tr><tr><td>non-user</td><td>2</td><td>6</td><td>8</td></tr><tr><td>All</td><td>321</td><td>308</td><td>629</td></tr></table><div>Precision: 0.9937694704049844 Recall: 0.5136876006441223</div></div>	Predicted	user	non-user	All	Actual				user	319	302	621	non-user	2	6	8	All	321	308	629	<div><div>Random Forest classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>435</td><td>186</td><td>621</td></tr><tr><td>non-user</td><td>6</td><td>2</td><td>8</td></tr><tr><td>All</td><td>441</td><td>188</td><td>629</td></tr></table><div>Precision: 0.9863945578231292 Recall: 0.7004830917874396</div></div>	Predicted	user	non-user	All	Actual				user	435	186	621	non-user	6	2	8	All	441	188	629
Predicted	user	non-user	All																																						
Actual																																									
user	319	302	621																																						
non-user	2	6	8																																						
All	321	308	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	435	186	621																																						
non-user	6	2	8																																						
All	441	188	629																																						
<div><div>SVM classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>469</td><td>152</td><td>621</td></tr><tr><td>non-user</td><td>7</td><td>1</td><td>8</td></tr><tr><td>All</td><td>476</td><td>153</td><td>629</td></tr></table><div>Precision: 0.9852941176470589 Recall: 0.7552334943639292</div></div>	Predicted	user	non-user	All	Actual				user	469	152	621	non-user	7	1	8	All	476	153	629	<div><div>Gradient Boosting classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>427</td><td>194</td><td>621</td></tr><tr><td>non-user</td><td>6</td><td>2</td><td>8</td></tr><tr><td>All</td><td>433</td><td>196</td><td>629</td></tr></table><div>Precision: 0.9861431870669746 Recall: 0.6876006441223832</div></div>	Predicted	user	non-user	All	Actual				user	427	194	621	non-user	6	2	8	All	433	196	629
Predicted	user	non-user	All																																						
Actual																																									
user	469	152	621																																						
non-user	7	1	8																																						
All	476	153	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	427	194	621																																						
non-user	6	2	8																																						
All	433	196	629																																						
<div><div>MLP classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>473</td><td>148</td><td>621</td></tr><tr><td>non-user</td><td>5</td><td>3</td><td>8</td></tr><tr><td>All</td><td>478</td><td>151</td><td>629</td></tr></table><div>Precision: 0.9895397489539749 Recall: 0.7616747181964574</div></div>	Predicted	user	non-user	All	Actual				user	473	148	621	non-user	5	3	8	All	478	151	629	<div><div>KNN classifier</div><table><tr><th>Predicted</th><th>user</th><th>non-user</th><th>All</th></tr><tr><th>Actual</th><td></td><td></td><td></td></tr><tr><td>user</td><td>448</td><td>173</td><td>621</td></tr><tr><td>non-user</td><td>5</td><td>3</td><td>8</td></tr><tr><td>All</td><td>453</td><td>176</td><td>629</td></tr></table><div>Precision: 0.9889624724061811 Recall: 0.7214170692431562</div></div>	Predicted	user	non-user	All	Actual				user	448	173	621	non-user	5	3	8	All	453	176	629
Predicted	user	non-user	All																																						
Actual																																									
user	473	148	621																																						
non-user	5	3	8																																						
All	478	151	629																																						
Predicted	user	non-user	All																																						
Actual																																									
user	448	173	621																																						
non-user	5	3	8																																						
All	453	176	629																																						

Mushrooms Dataset

Decision Tree classifier

Predicted	user	non-user	All
Actual			
user	388	233	621
non-user	3	5	8
All	391	238	629

Precision: 0.9923273657289002
Recall: 0.6247987117552335

Random Forest classifier

Predicted	user	non-user	All
Actual			
user	435	186	621
non-user	6	2	8
All	441	188	629

Precision: 0.9863945578231292
Recall: 0.7004830917874396

SVM classifier

Predicted	user	non-user	All
Actual			
user	469	152	621
non-user	7	1	8
All	476	153	629

Precision: 0.9852941176470589
Recall: 0.7552334943639292

Gradient Boosting classifier

Predicted	user	non-user	All
Actual			
user	427	194	621
non-user	6	2	8
All	433	196	629

Precision: 0.9861431870669746
Recall: 0.9861431870669746

MLP classifier

Predicted	user	non-user	All
Actual			
user	474	147	621
non-user	5	3	8
All	479	150	629



Precision: 0.9895615866388309
Recall: 0.7632850241545893

KNN classifier

Predicted	user	non-user	All
Actual			
user	448	173	621
non-user	5	3	8
All	453	176	629

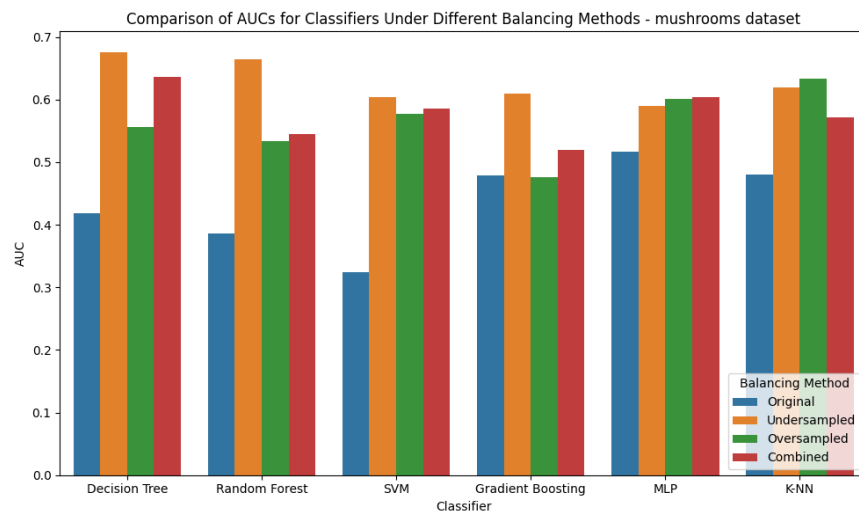
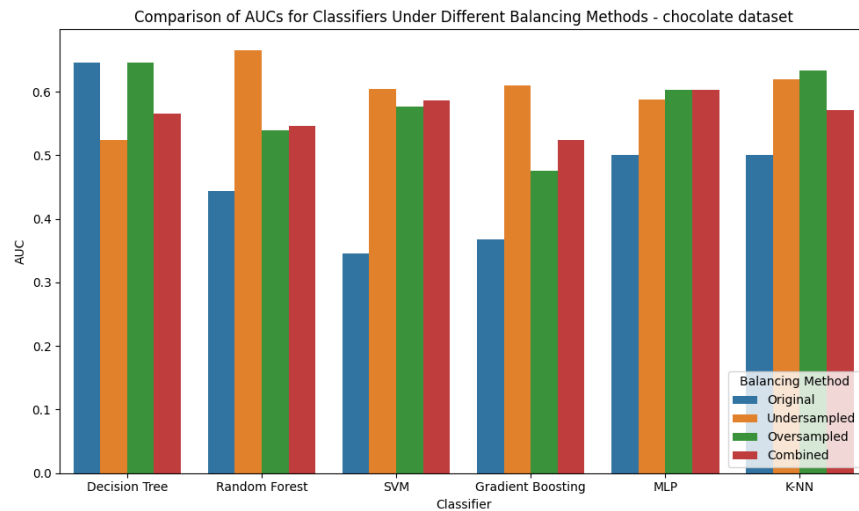
Precision: 0.9889624724061811
Recall: 0.7214170692431562

ROC curve as a Summary



Further Analysis

Based on the above estimations, essentially the ROC curves, we have the AUCs which compares the performance of different classifiers under different sampling methods. Therefore, they are summarized in the following bar plots.



Lesson Learnt

Processes in the Pipeline

The pipeline begins with data preprocessing, followed by feature extraction. Six classifiers are then instantiated with custom parameters. *Random Search* was used for hyperparameter optimization. After splitting the datasets into training and test sets, classifiers were trained and evaluated using precision, recall, and AUC.

Analysis from the Evaluation Metrics

The primary analysis is mainly based on the *precision*, *recall*, and *AUC* scores from confusion matrices. In the original classification (without resampling), the Decision Tree achieved the highest scores across both datasets. With *under-sampling*, the Random Forest showed the highest precision in both datasets, while the Decision Tree had the highest recall and AUC in the chocolate dataset. For *oversampling*, KNN gave the highest precision and AUC in the chocolate dataset, and the SVM had the highest recall. Using *combined sampling*, the Decision Tree excelled in precision and AUC in both datasets, while MLP and Gradient Boosting had the highest recall. Overall, the **Decision Tree** is the best classifier which is also suggested in the bar plot regarding AUC comparisons across different sampling methods.

Issue in the Datasets

Why different sampling methods suggests the best classifier differently is because of the *class imbalance* issue. In the Chocolate dataset, we can clearly see that the number of “user” category sample (i.e., 621) is far more than that of “non-user” category samples (i.e., 8). In the Mushrooms dataset, we can also clearly see that the number of magic mushroom non-users (i.e., 386) is slightly more than that of users (i.e., 243). The imbalance of classes obviously leads to a biased conclusion. In other words, all techniques generally suggest **decision tree** the best classifier.

Addressing the Issue

Regarding the *class imbalance*, *Random Under Sampler* is used to under-sample the majority classes in both datasets; *SMOTE* is used to over-sample the minority classes; and a combination of techniques is used to make a fair sampling approach to the data. The advantage of that is obviously balancing the number of samples in each class. However, this may lead to other potential problems. For instance, after under-sampling the dataset, some data from the majority class are trimmed. This causes data loss. On the other hand, after over-sampling the minority class, some unnecessary data are added. Affirmatively, it reduces biases in the results, but it increases the variance of data.

Conclusion

According to the notion of “No Free Lunch” principle, no single algorithm is always the most accurate one. Therefore, the analysis here also considers using a combination of sampling methods to address the *class imbalance* issue fairly. The AUC comparison graphs suggests that:

The combination of methods has nearly an average AUC between both sampling methods; and,

By using any sampling techniques (including the combination of methods), it generally gives a higher AUC score from any classifier based on any dataset.

In conclusion, by using a sampling technique reduces biases from the class imbalance issue; and by using a combination of sampling techniques, a fair conclusion is drawn.